

MUSIC GENRE CLASSIFICATION

Corvin Idler | 23.6.2022

CONTENTS



1

Introduction

2

Features and Pre-Processing

3

Algorithm

4

Results

THE FREE MUSIC ARCHIVE (FMA)

The FMA consists of A LOT! of Creative Commons licensed audio.
<https://freemusicarchive.org/home>

Luckily somebody extracted a lot of audio features already:
<https://github.com/mdeff/fma>

FMA in numbers

106,574 tracks from 16,341 artists and 14,854 albums, arranged in a hierarchical taxonomy of 161 genres grouped together into 16 top level genres

Question

Can we do some music genre classification with these audio features?

Answer

Yes... well.. I guess it depends :)

Detailed results

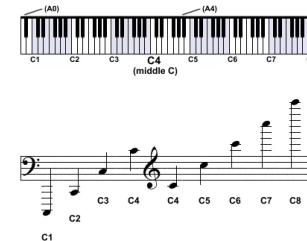
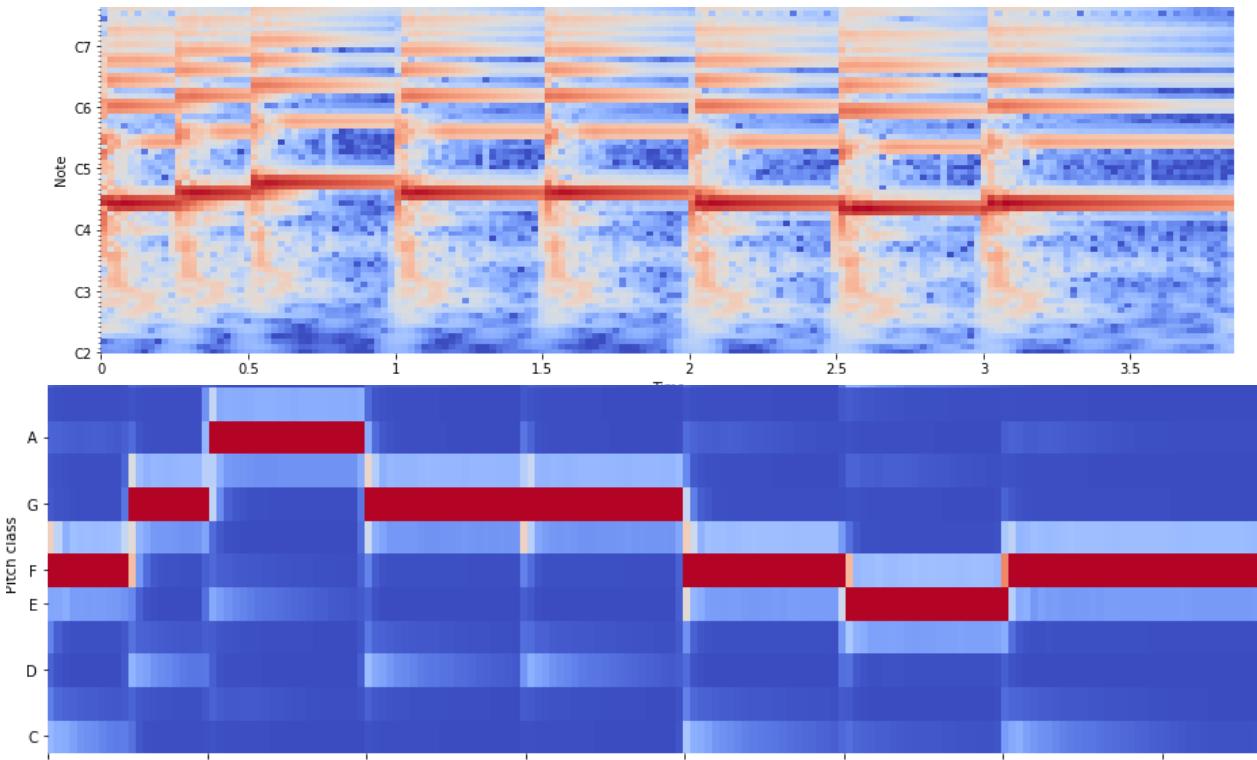
Fasten your seatbelt and stay tuned :)

AUDIO FEATURES PRIMER

A spectrogram (e.g. Fourier Transformation) is turned into a chromatogram when “broken down” by the 12 pitch classes (C, C#, D, D#, E, F, F#, G, G#, A, A#, and B)

Each of these pitch classes exist several times (an 88-key piano has seven octaves)

The chromatogram tries to capture the energy for all octaves but grouped into pitch classes.



Octaves

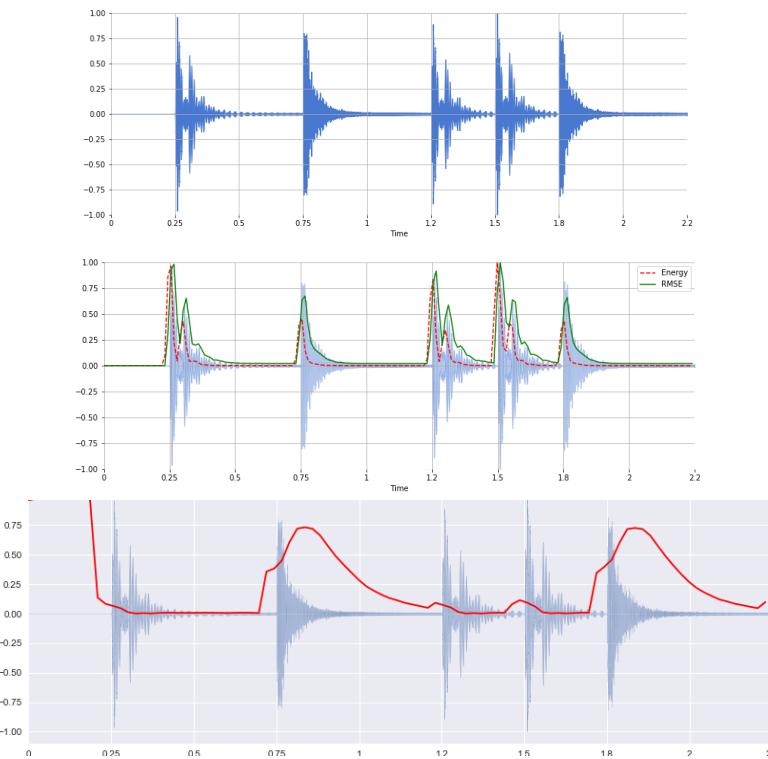
AUDIO FEATURES CONTINUED

There are also slightly less sophisticated feature one can calculate

Zero crossing rate, Root Mean Square Energy, Spectral Centroid, Spectral Bandwidth, etc.

These features need to cover the whole song. So if calculated for many time slices we get a distribution not a scalar!

-> kurtosis, max, mean, median, min, skew and standard deviation were calculated to create scalar features for a whole song



Have a look at

<https://github.com/librosa/librosa>

<https://github.com/mdeff/fma>

<https://musicinformationretrieval.com/>

EDA / PRE-PROCESSING

Genre	count
NaN	56976
Rock	14182
Experimental	10608
Electronic	9372
Hip-Hop	3552
Folk	2803
Pop	2332
Instrumental	2079
International	1389
Classical	1230
Jazz	571
Old-Time/Historic	554
Spoken	423
Country	194
Soul-RnB	175
Blues	110
Easy Listening	24

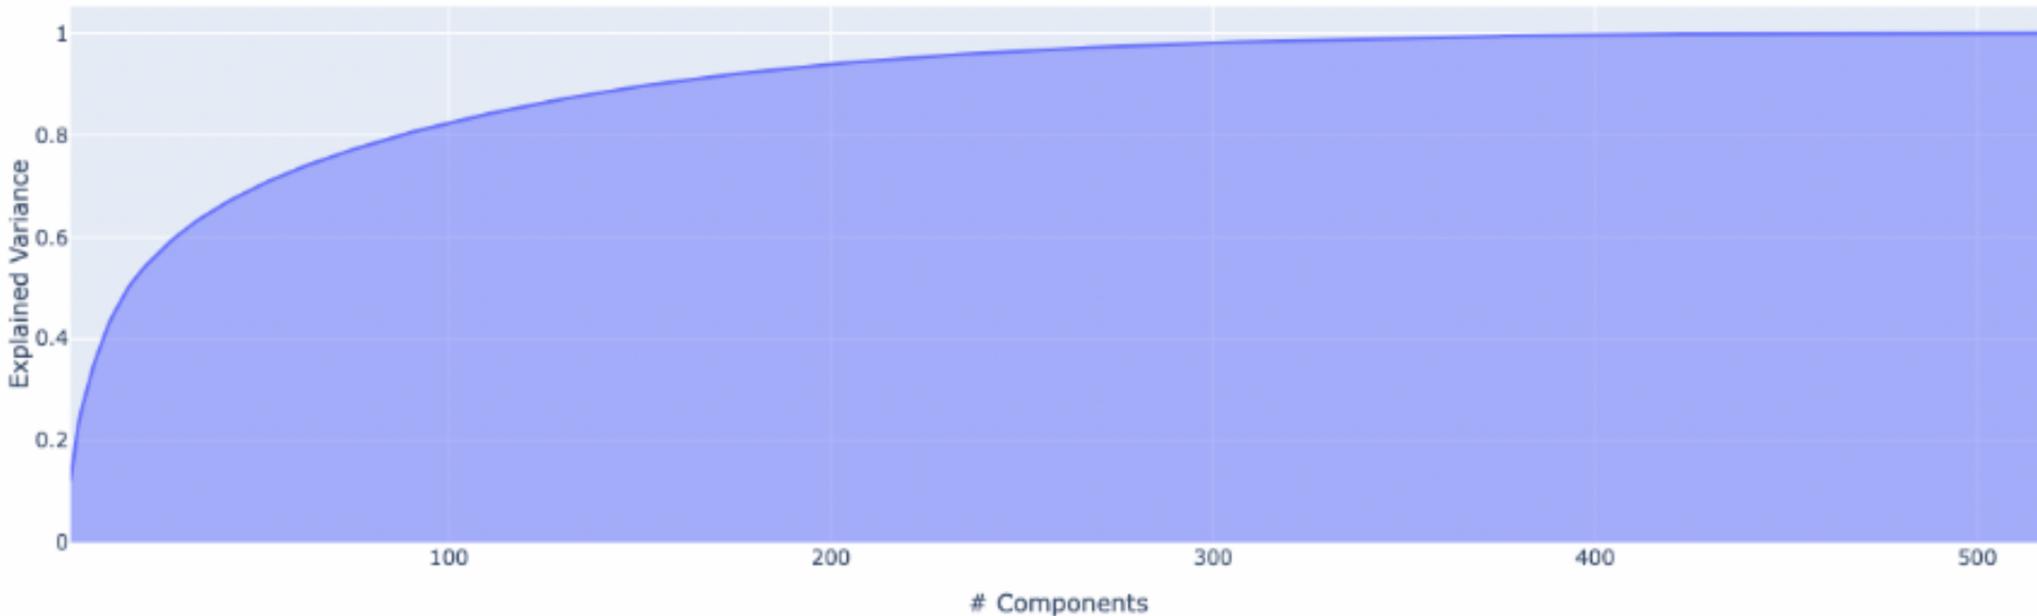
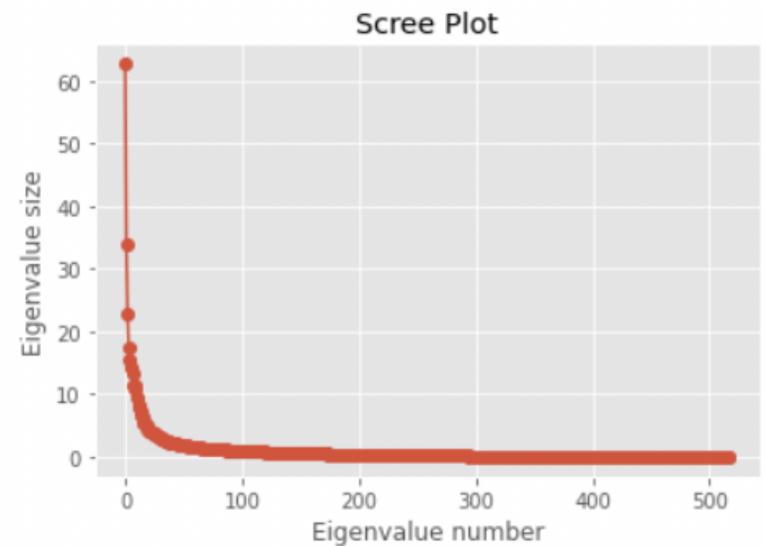
Quite some NaNs for the Genres, and not all features particularly “useful”

feature	mean	std	min	25%	50%	75%	max
chroma stft.max.12	0.99	0.003	0.79	1.0	1.0	1.0	1.0
chroma stft.max.10	0.99	0.003	0.51	1.0	1.0	1.0	1.0
chroma stft.max.01	0.99	0.004	0.66	1.0	1.0	1.0	1.0
chroma stft.max.05	0.99	0.004	0.68	1.0	1.0	1.0	1.0
chroma stft.max.03	0.99	0.004	0.54	1.0	1.0	1.0	1.0
chroma stft.max.04	0.99	0.004	0.49	1.0	1.0	1.0	1.0
chroma stft.max.06	0.99	0.004	0.66	1.0	1.0	1.0	1.0
chroma stft.max.02	0.99	0.005	0.72	1.0	1.0	1.0	1.0
chroma stft.max.11	0.99	0.005	0.58	1.0	1.0	1.0	1.0
chroma stft.max.08	0.99	0.006	0.53	1.0	1.0	1.0	1.0
chroma stft.max.09	0.99	0.006	0.58	1.0	1.0	1.0	1.0
chroma stft.max.07	0.99	0.007	0.48	1.0	1.0	1.0	1.0
chroma cqt.max.01	0.99	0.007	0.49	1.0	1.0	1.0	1.0
chroma cqt.max.03	0.99	0.009	0.32	1.0	1.0	1.0	1.0
chroma cqt.max.04	0.99	0.009	0.31	1.0	1.0	1.0	1.0

...

PCA RESULTS

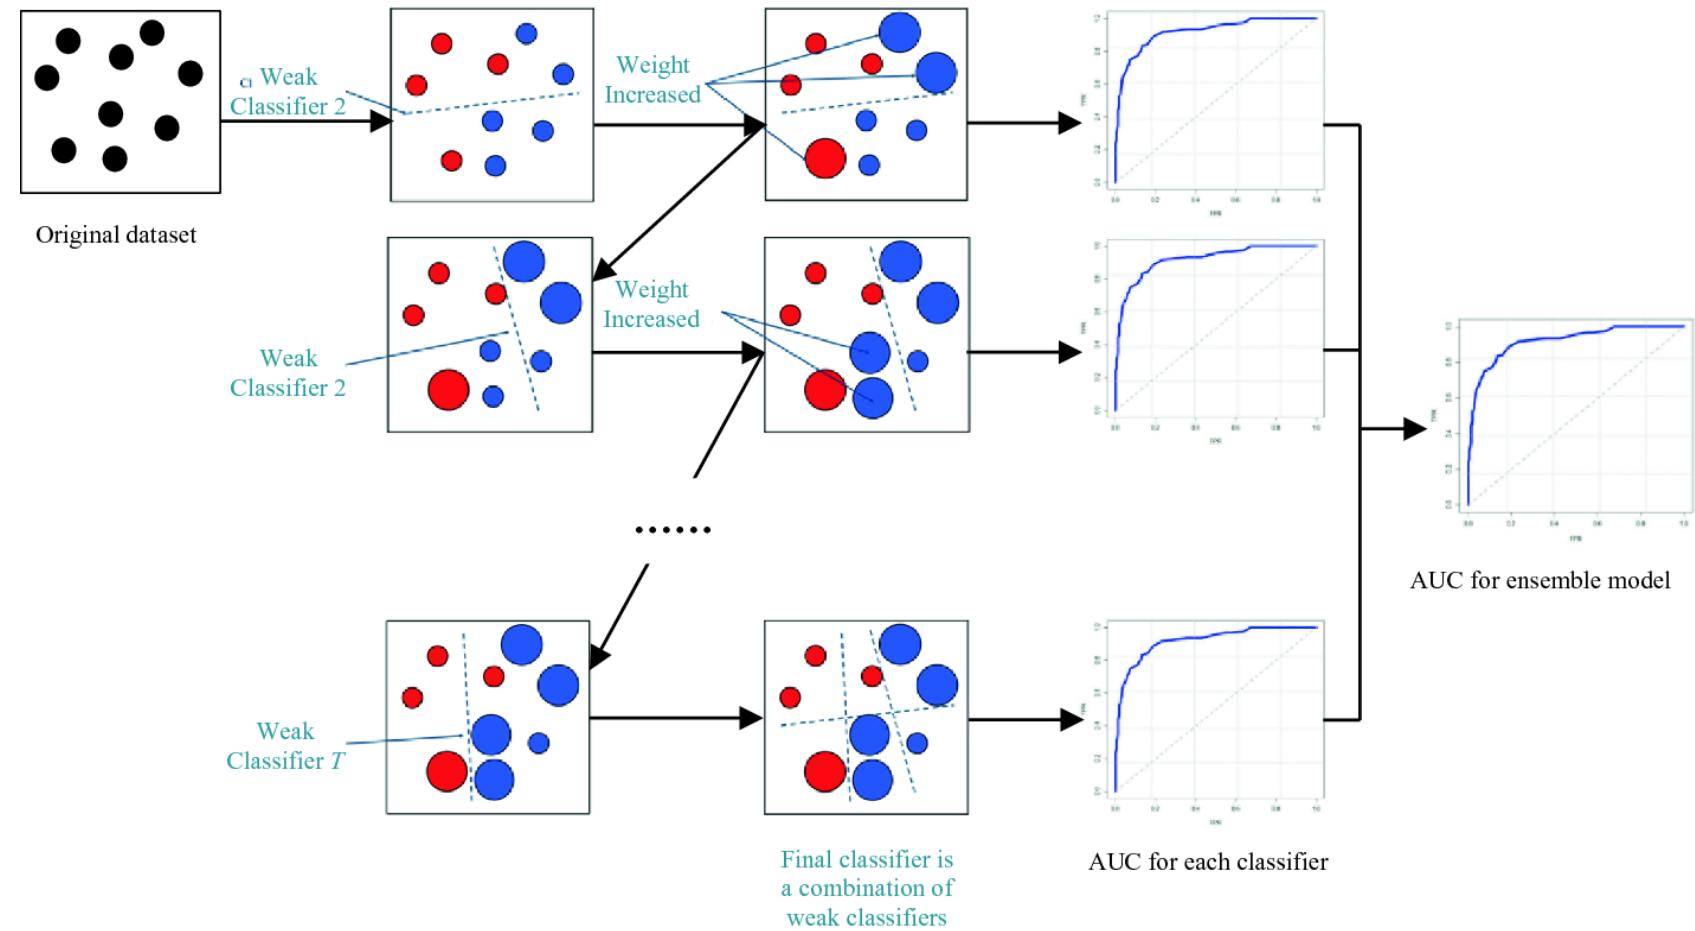
"elbow" in Scree plot around 20-30 factors
but ca. 100 factors needed for 80% variance...
so a loooooong tail...



GRADIENT BOOSTED DECISION TREES

LightGBM

- multi class
- Gradient-based One-Side Sampling
- Mutual exclusion feature binding
- runs on "anything" when installed





NOT MUCH TO BE SEEN!?

A CLOSER LOOK:

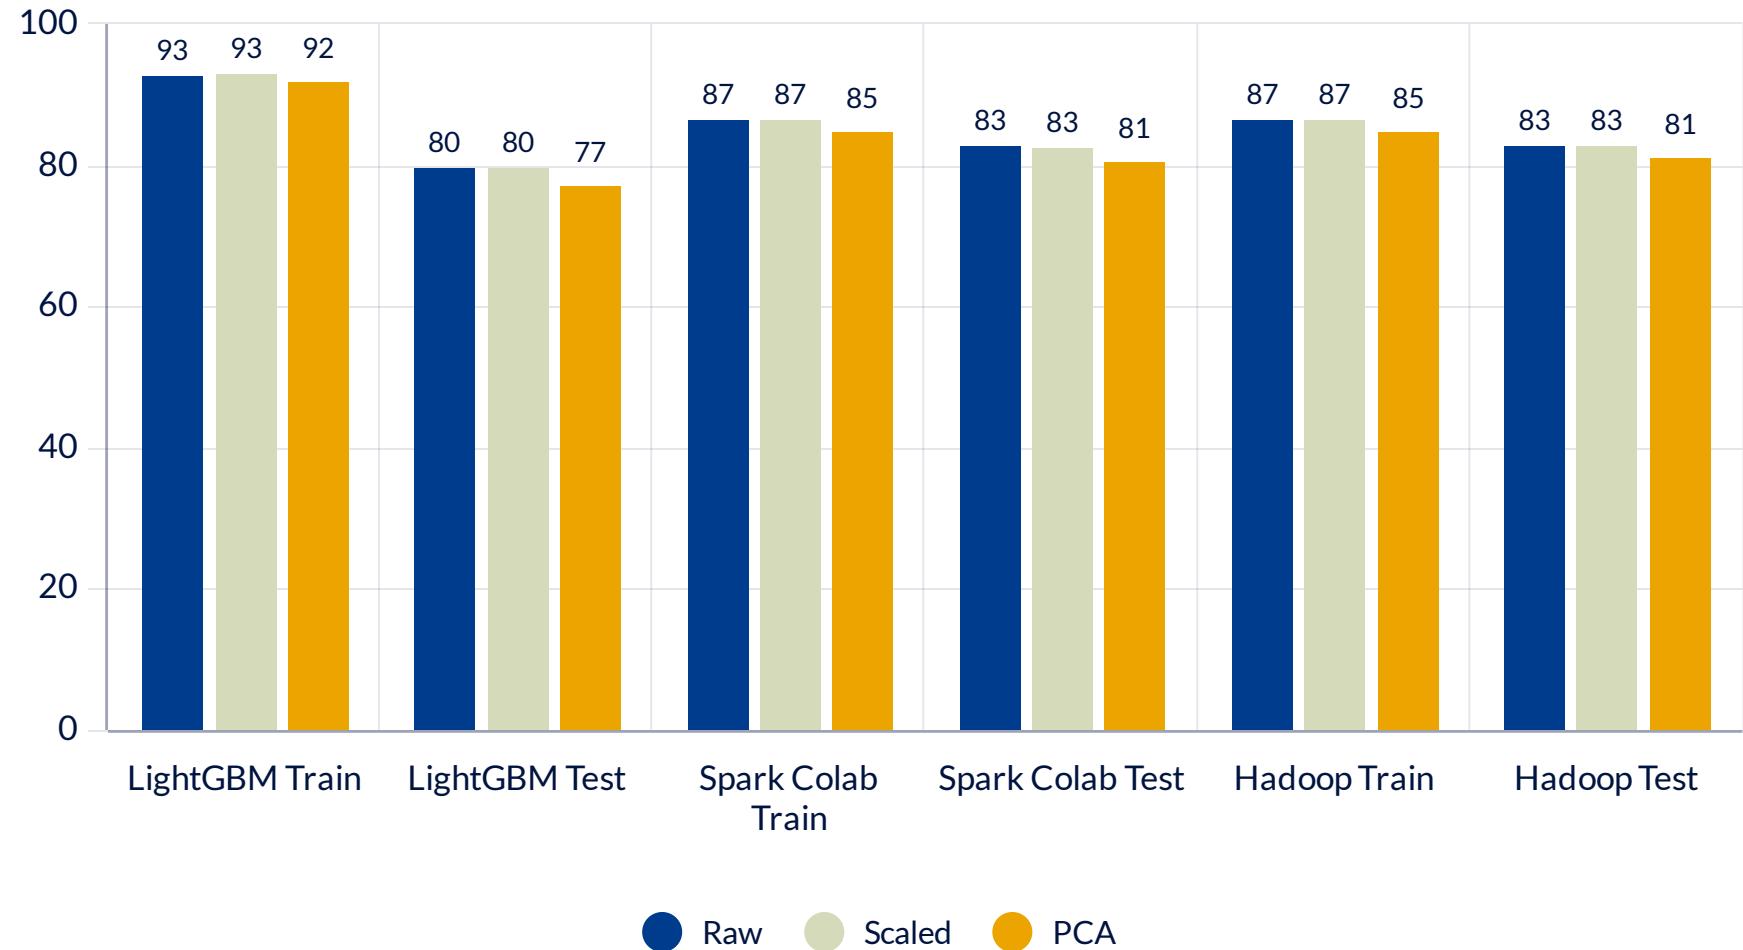
LightGBM (100 trees)
overfit more than
Spark (20 trees)

Raw or scaled doesn't
matter much for a tree

PCA is always hurting.
Makes sense for GBM

RESULTS - ACCURACY

Train vs. Test and Raw vs. Scale vs PCA



RESULTS - CONFUSION MATRIX

LightGBM results

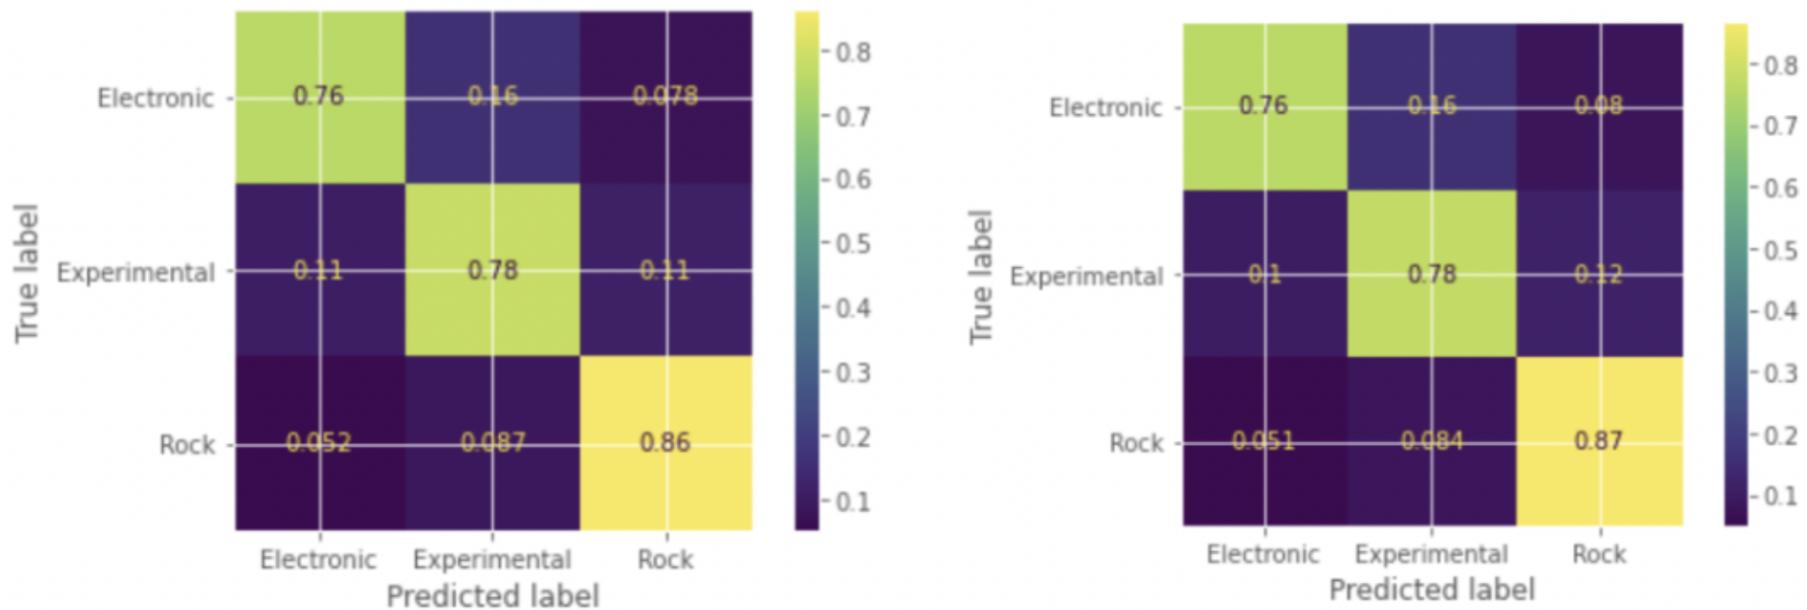


Figure 5: Confusion matrices for raw features (left) vs. scaled/normalized (right)



WHAT ELSE?

A CLOSER LOOK:

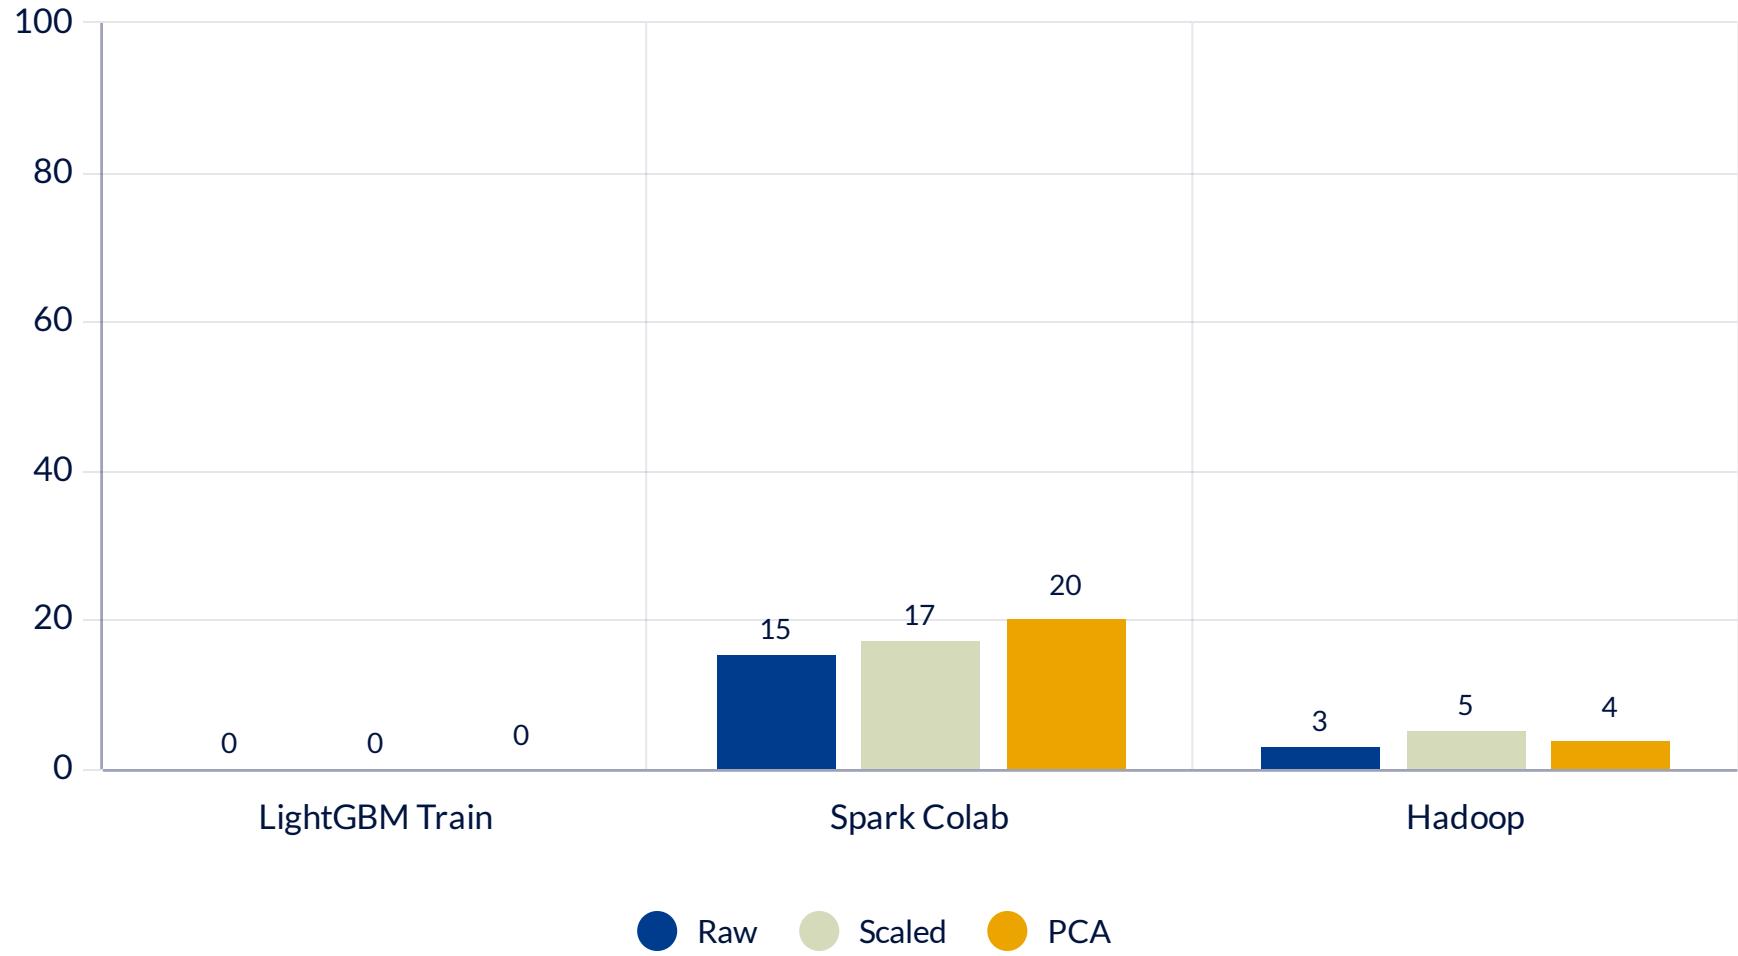
LightGBM (100 trees)
but waaay faster

Colab is free but
certainly not
overperforming...

Hadoop takes still
quite a bit.. Maybe
random forest better
for distributed
computing

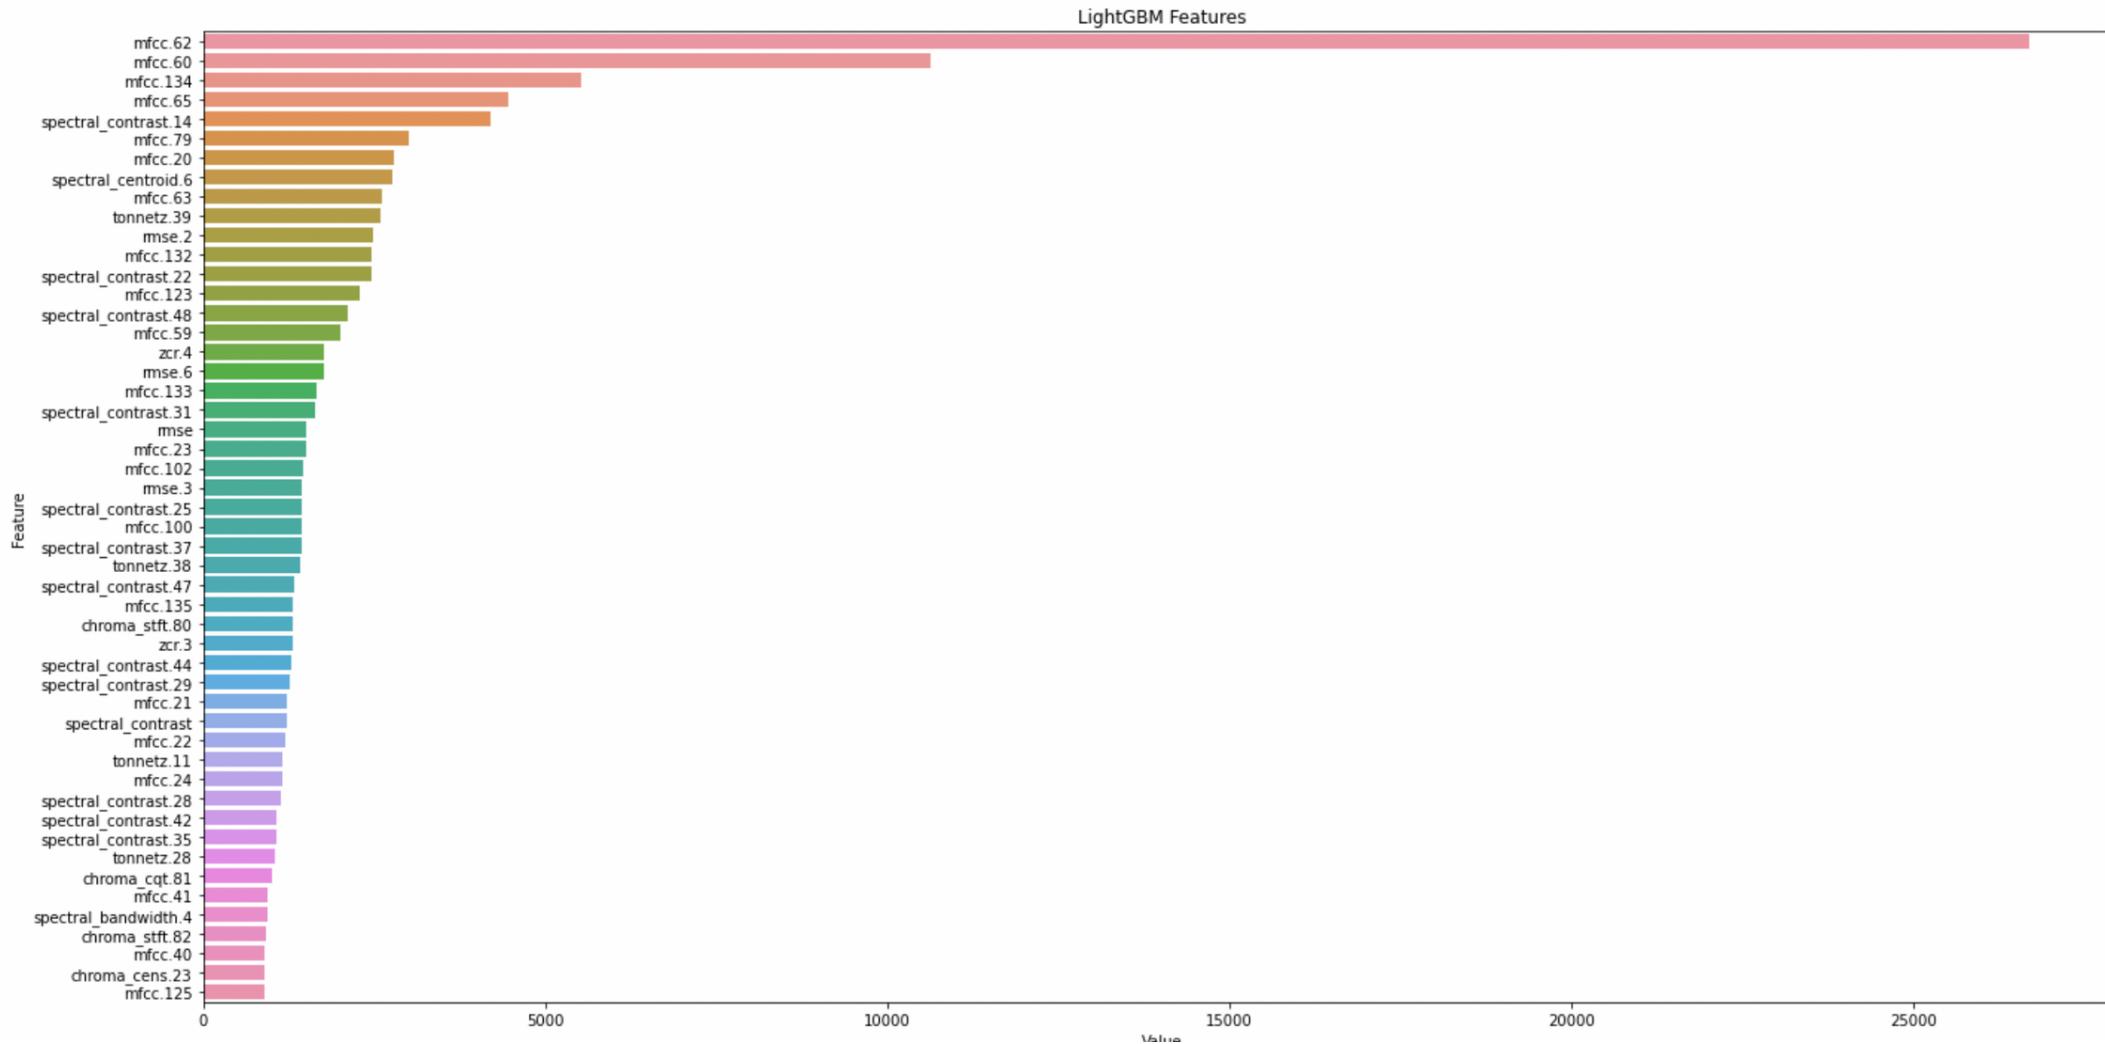
RESULTS RUN TIME

timing matters... values in minutes



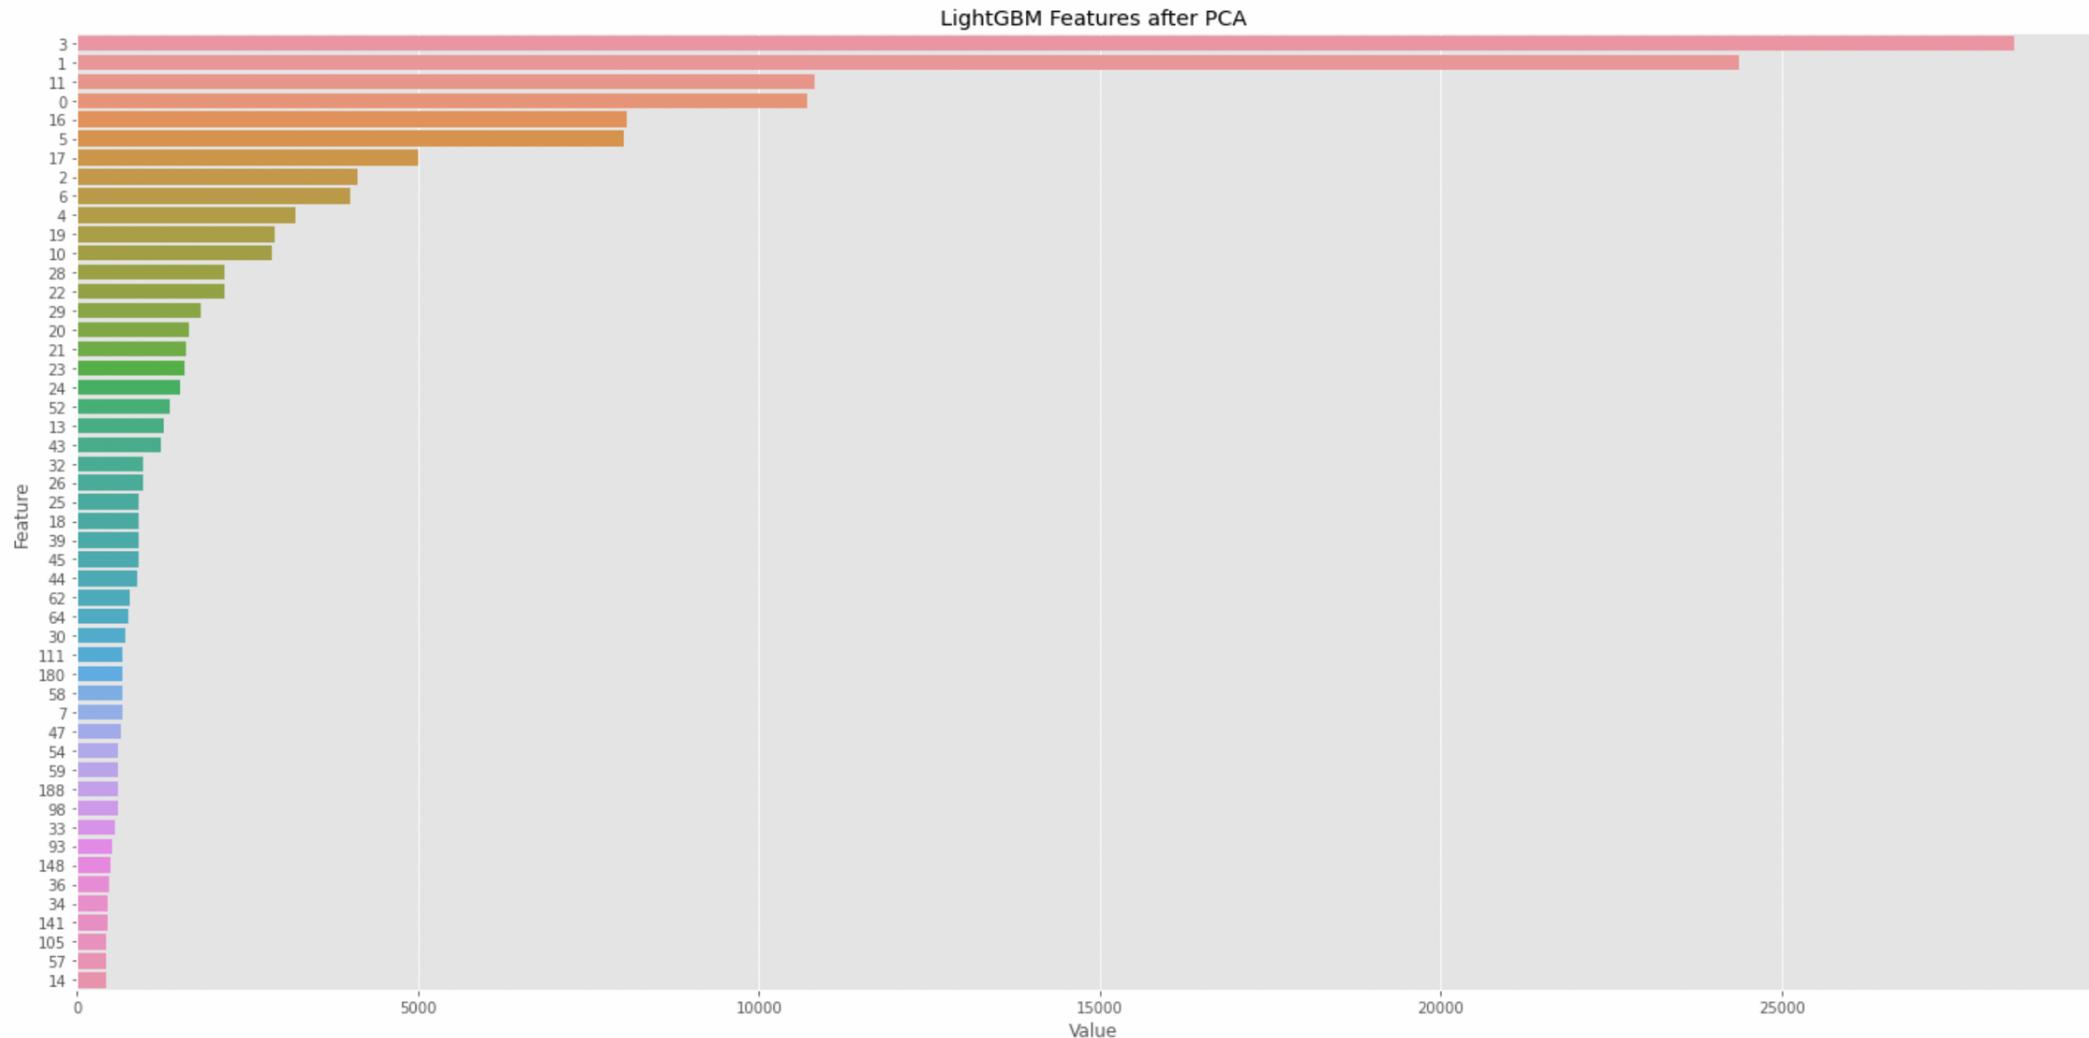
RESULTS FEATURE IMPORTANCE

Feature importance - raw features



RESULTS FEATURE IMPORTANCE PCA

Feature importance - principal components



CONCLUSIONS



1

a highly optimized algorithm on a desktop might actually out-compete an “aged” implementation on a cluster in terms of runtime



2

Point #2
A cluster might come in handy when the task is highly parallelizable. E.g. random forest vs. GBDT or hyper-parameter optimization vs. single model.



3

The principal components with the largest eigenvalue might not at all be the most important feature.
Succinct findings + recommendations

QUESTIONS?